

Enterprise RAG Platform — Complete Page Breakdown

PAGE 1 — Dashboard (Home)

****Route:**** `/dashboard`

****Kaam:**** Poore system ka ek nazar mein overview

****Section 1 — Top Navbar****

- Logo (left side)
- Navigation links — Dashboard, Documents, Chat, Evaluation, Search Lab, Feedback, Settings
- Notification bell icon — system alerts
- User avatar + dropdown — profile, logout

****Section 2 — Stats Bar (4 cards)****

- Total Documents uploaded — count
- Queries Today — how many RAG queries ran
- Average Faithfulness Score — RAGAS score today
- Active Chat Sessions — live sessions count

****Section 3 — Recent Queries Feed****

- Last 10 queries list
- Each row shows — question text, timestamp, faithfulness score, pass/fail badge
- Click karo toh chat page pe jao us query ke saath

****Section 4 — Quick Actions****

- Upload New PDF button
- Start New Chat button
- Run Evaluation button
- View Reports button

****Section 5 — System Health Panel****

- Vector DB status — green/red dot
- Redis Cache status — connected or not
- Embedding Model status — loaded or not
- API response latency — ms meter
- Last ingestion time — "5 mins ago"

****Section 6 — Top Documents****

- Most queried documents list
- Document name, query count, last accessed

PAGE 2 — Document Manager

****Route:**** `/documents`

****Kaam:**** PDF upload karo, chunks banao, documents manage karo

****Section 1 — Upload Zone****

- Drag & drop area — large box with upload icon
- Browse files button
- Supported formats label — PDF, TXT, DOCX
- Bulk upload support — multiple files ek saath
- Upload progress bar — file by file

****Section 2 — Chunking Strategy Panel****

- Strategy dropdown —
 - Fixed Size (chunk size + overlap slider)
 - Recursive Character (separators list)
 - Semantic Chunking (similarity threshold)
 - Sliding Window (window size + step)
- Chunk preview — shows how text will be split
- Estimated chunks count — live update

****Section 3 — Metadata Editor****

- Document title field
- Category dropdown — Finance, Legal, HR, Technical, Other
- Tags input — add multiple tags
- Source URL field — optional
- Date field — document date
- Description textarea

****Section 4 — Document List Table****

- Columns — Name, Size, Pages, Chunks, Status, Upload Date, Actions
- Status badges — Processing, Embedded, Failed, Ready
- Search bar — filter by name
- Filter dropdown — by category, by status, by date
- Sort — by date, by name, by size
- Pagination — 10/25/50 per page

****Section 5 — Document Detail Drawer (right side)****

- Opens on clicking any document
- Shows — full metadata
- Chunk viewer — all chunks listed with chunk number
- Embedding status per chunk

- Delete document button
- Re-process button — re-chunk with new strategy

****Section 6 — Processing Pipeline Indicator****

- Live status steps —
 - Step 1 — Uploading ☐
 - Step 2 — Parsing Text ☐
 - Step 3 — Chunking ☐
 - Step 4 — Generating Embeddings ☐
 - Step 5 — Storing in Vector DB ☐
 - Step 6 — Ready ☐

PAGE 3 — Chat / Query Interface

****Route:**** `/chat`

****Kaam:**** Documents se sawaal poochho, citations ke saath jawab lo

****Section 1 — Left Sidebar****

- New Chat button — top
- Chat history list — previous conversations
- Each item — first message preview + date
- Search chats — search bar
- Delete chat option — hover pe aata hai

****Section 2 — Top Query Settings Bar****

- Model selector — GPT-4, Claude, Gemini dropdown
- Temperature slider — 0.0 to 1.0
- Top-K slider — how many chunks retrieve (1-20)
- Search mode toggle — Vector / Hybrid
- Reranking toggle — on/off switch
- Document filter — select specific documents to query

****Section 3 — Main Chat Area****

- User message bubble — right side, blue
- AI response bubble — left side, white
- Each AI response has —
 - Answer text
 - "Sources used" expandable section
 - Confidence score badge
 - Copy button
 - Thumbs up / Thumbs down feedback

- Typing indicator — "RAG is thinking..."

- Scroll to bottom button

****Section 4 — Citation Panel (right side)****

- Opens when AI responds

- Shows retrieved chunks used

- Each chunk shows —

- Document name

- Page number

- Chunk text (highlighted)

- Similarity score

- Rank number (after reranking)

- Toggle — show all chunks / show top 3 only

****Section 5 — Query Rewriting Preview****

- Expandable panel below input

- Shows — Original Query

- Shows — Rewritten Query (LLM rewritten)

- Shows — Multi-queries generated (3 sub-questions)

- Toggle on/off

****Section 6 — Input Area****

- Text input box — multiline

- Send button

- Voice input button — mic icon

- Attach document button — query specific doc

- Character / token counter

- Conversation memory indicator — "8 messages in memory"

- Clear memory button

PAGE 4 — Evaluation Dashboard

****Route:**** `/evaluation`

****Kaam:**** RAGAS metrics se RAG pipeline ki quality measure karo

****Section 1 — Score Overview Cards (4 cards)****

- Faithfulness — score 0 to 1, color coded

- Context Recall — score 0 to 1

- Context Precision — score 0 to 1

- Answer Relevancy — score 0 to 1

- Each card shows — current score, previous score, up/down arrow

****Section 2 — Run Evaluation Panel****

- Dataset selector — upload test CSV or use existing
- Metric checkboxes — select which metrics to run
- LLM selector — which model to use for evaluation
- Run Evaluation button
- Estimated time label — "~2 mins for 50 queries"
- Progress bar — live while running

****Section 3 — Score Trend Chart****

- Line graph — X axis dates, Y axis scores
- 4 lines — one per metric, different colors
- Date range filter — Last 7 days / 30 days / custom
- Hover tooltip — exact score on that date

****Section 4 — Per-Query Evaluation Table****

- Columns —
 - Question
 - Generated Answer
 - Ground Truth
 - Faithfulness score
 - Context Recall score
 - Context Precision score
 - Answer Relevancy score
 - Overall Pass / Fail
- Row click — expand full detail
- Filter — show only failed queries
- Sort by any score column

****Section 5 — Failure Analysis Panel****

- Auto-grouped failures —
 - Hallucination issues (low faithfulness)
 - Missing context (low recall)
 - Irrelevant retrieval (low precision)
 - Off-topic answers (low relevancy)
- Count per failure type
- Sample queries for each type

****Section 6 — Export Section****

- Download CSV — full evaluation results
- Download PDF Report — formatted report
- Share link — shareable evaluation report URL

PAGE 5 — Search & Retrieval Lab

****Route:**** `/search`

****Kaam:**** Retrieval pipeline test karo, hybrid search experiment karo

****Section 1 — Search Input****

- Large search box — center
- Search button
- Example queries — 3 clickable sample queries

****Section 2 — Search Configuration Panel****

- Search Mode —
 - Vector Search only
 - BM25 Keyword only
 - Hybrid (BM25 + Vector) — weight slider
- Top-K selector — how many results (1 to 20)
- Reranking toggle — on/off
- Reranker model selector — Cohere, BGE, Cross-encoder

****Section 3 — Query Processing Preview****

- Original query display
- Query rewriting result — how LLM rewrote it
- Multi-query expansion — 3 generated sub-queries
- Embeddings vector preview — first 10 dimensions shown

****Section 4 — Metadata Filter Sidebar****

- Filter by Category — checkboxes
- Filter by Tags — multi-select
- Filter by Date range — date picker
- Filter by Document — select specific files
- Apply filters button
- Reset filters button

****Section 5 — Results Panel****

- Each result card shows —
 - Rank number (1, 2, 3...)
 - Document name + page
 - Chunk text — highlighted matching words
 - Vector similarity score
 - BM25 score (if hybrid)
 - Final reranked score
 - Expand full chunk button

- Side by side comparison — Vector results vs Hybrid results
- Result count — "Showing 8 of 47 chunks"

****Section 6 — Retrieval Analytics****

- Query latency — ms
- Vector search time vs BM25 time breakdown
- Chunks scanned count
- Reranking time
- Total pipeline time

PAGE 6 — Feedback & Analytics

****Route:**** `/feedback`

****Kaam:**** User feedback collect karo, patterns samjho, bad responses fix karo

****Section 1 — Summary Cards (3 cards)****

- Total Feedback collected
- Positive responses % — thumbs up
- Negative responses % — thumbs down
- Needs Review count — flagged responses

****Section 2 — Feedback Table****

- Columns —
 - Query text
 - AI Response preview
 - User rating — thumbs up/down
 - User comment (if given)
 - Timestamp
 - Document used
 - Action — Mark reviewed, Delete
- Filter — positive only / negative only / unreviewed
- Search — search by query text

****Section 3 — Bad Response Review Panel****

- List of thumbs-down responses
- Each shows —
 - Original question
 - Bad AI answer
 - Retrieved chunks that were used
 - What went wrong — tag (hallucination, wrong doc, off-topic)
 - Mark as Fixed button

- Add to test dataset button

****Section 4 — Analytics Charts****

- Queries per day — bar chart
- Rating trend — positive vs negative over time line chart
- Top failure topics — pie chart
- Most queried documents — horizontal bar chart
- Peak usage hours — heatmap

****Section 5 — Test Dataset Builder****

- Add feedback items to evaluation dataset
- Question + Ground truth editor
- Export as RAGAS-ready CSV
- Dataset list — saved datasets

****Section 6 — Export****

- Export all feedback — CSV
- Export analytics report — PDF
- Date range selector for export

PAGE 7 — Settings

****Route:**** `/settings`

****Kaam:**** Poore system ka configuration ek jagah

****Section 1 — LLM Configuration****

- Provider selector — OpenAI, Anthropic, Google, Local
- Model name input — gpt-4o, claude-3-5, gemini-pro
- API Key field — masked input
- Temperature default — slider
- Max tokens — number input
- Test connection button

****Section 2 — Embedding Model****

- Model selector — text-embedding-3-small, ada-002, BGE, custom
- Dimensions display — auto-filled
- Batch size — number input
- Test embedding button

****Section 3 — Vector Database****

- DB selector — Qdrant / pgvector
- Host URL input
- Port input

- Collection name input
- Test connection button
- Re-index all documents button

****Section 4 — Chunking Defaults****

- Default strategy dropdown
- Default chunk size — slider (100-2000)
- Default overlap — slider (0-500)
- Apply to new uploads only toggle
- Re-process all documents with new settings button

****Section 5 — Redis Configuration****

- Host input
- Port input
- Password field — masked
- TTL (cache expiry) — slider in minutes
- Cache toggle — enable/disable
- Flush cache button

****Section 6 — Observability****

- LangSmith toggle — on/off
- LangSmith API key field
- Langfuse toggle — on/off
- Langfuse API key field
- Trace all queries toggle
- View traces button — opens LangSmith/Langfuse

****Section 7 — User Profile****

- Name field
- Email field
- Password change — old + new + confirm
- Profile picture upload
- Save changes button

Quick Summary

Page	Sections	Main Feature
Dashboard	6	System overview
Document Manager	6	Upload + chunking
Chat Interface	6	RAG queries + citations

Evaluation	6	RAGAS scoring
Search Lab	6	Retrieval testing
Feedback	6	Quality improvement
Settings	7	Full configuration

Kaunsa page ka mockup pehle banao? Dashboard ya Chat Interface? ☐